

Neo4j

Graph Databases

Agenda

- Overview
- Main Features
- Neo4j Architecture
- Neo4j Storage Strategy
- Neo4j based Application Architecture
- Benchmarks & Datasets

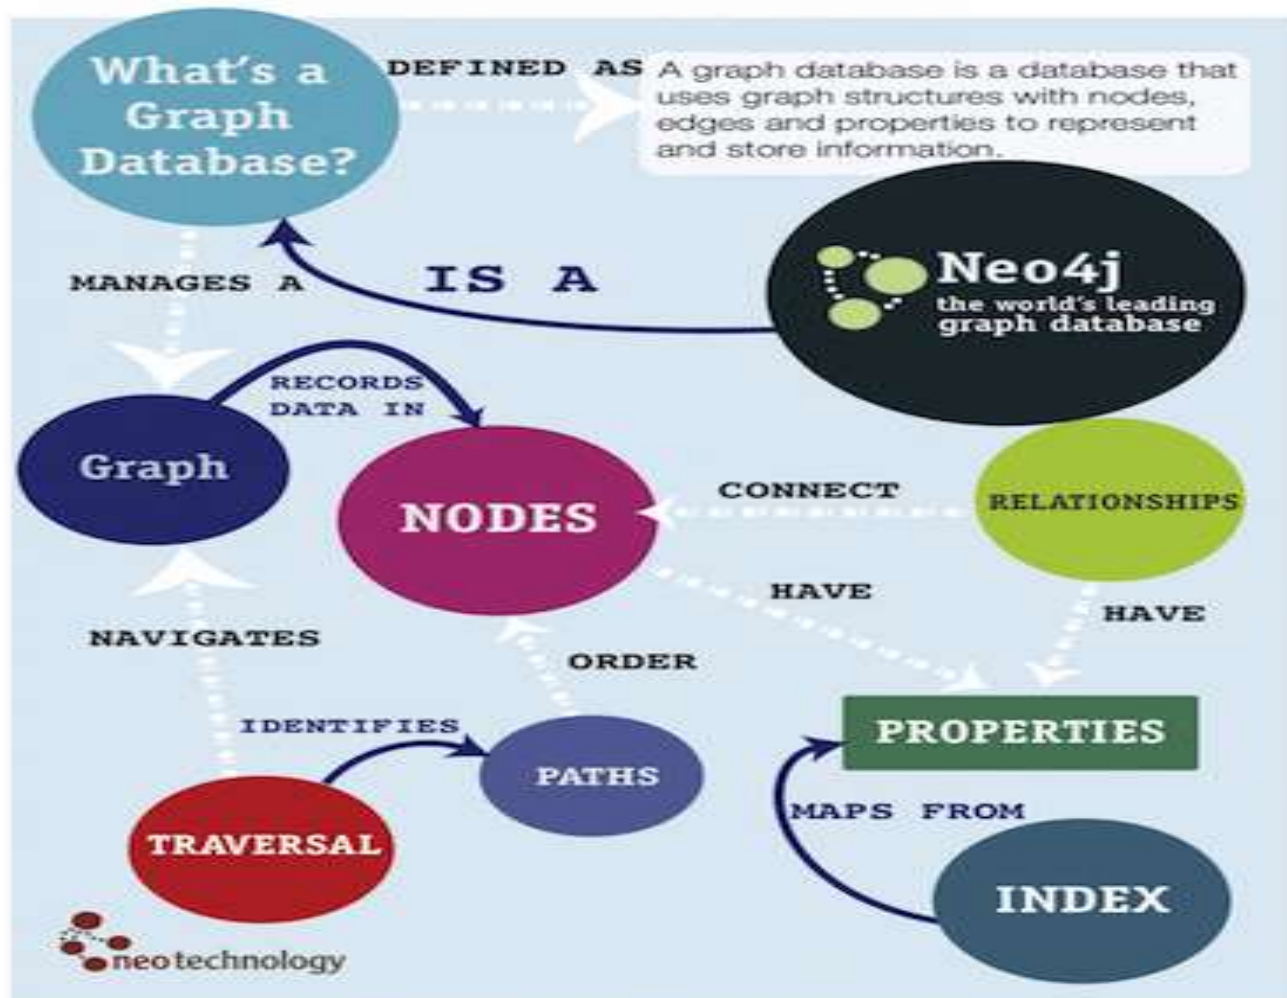
What is Graph Database?

- A graph database stores data in a graph, the most generic of data structures, capable of elegantly representing any kind of data in a highly accessible way.

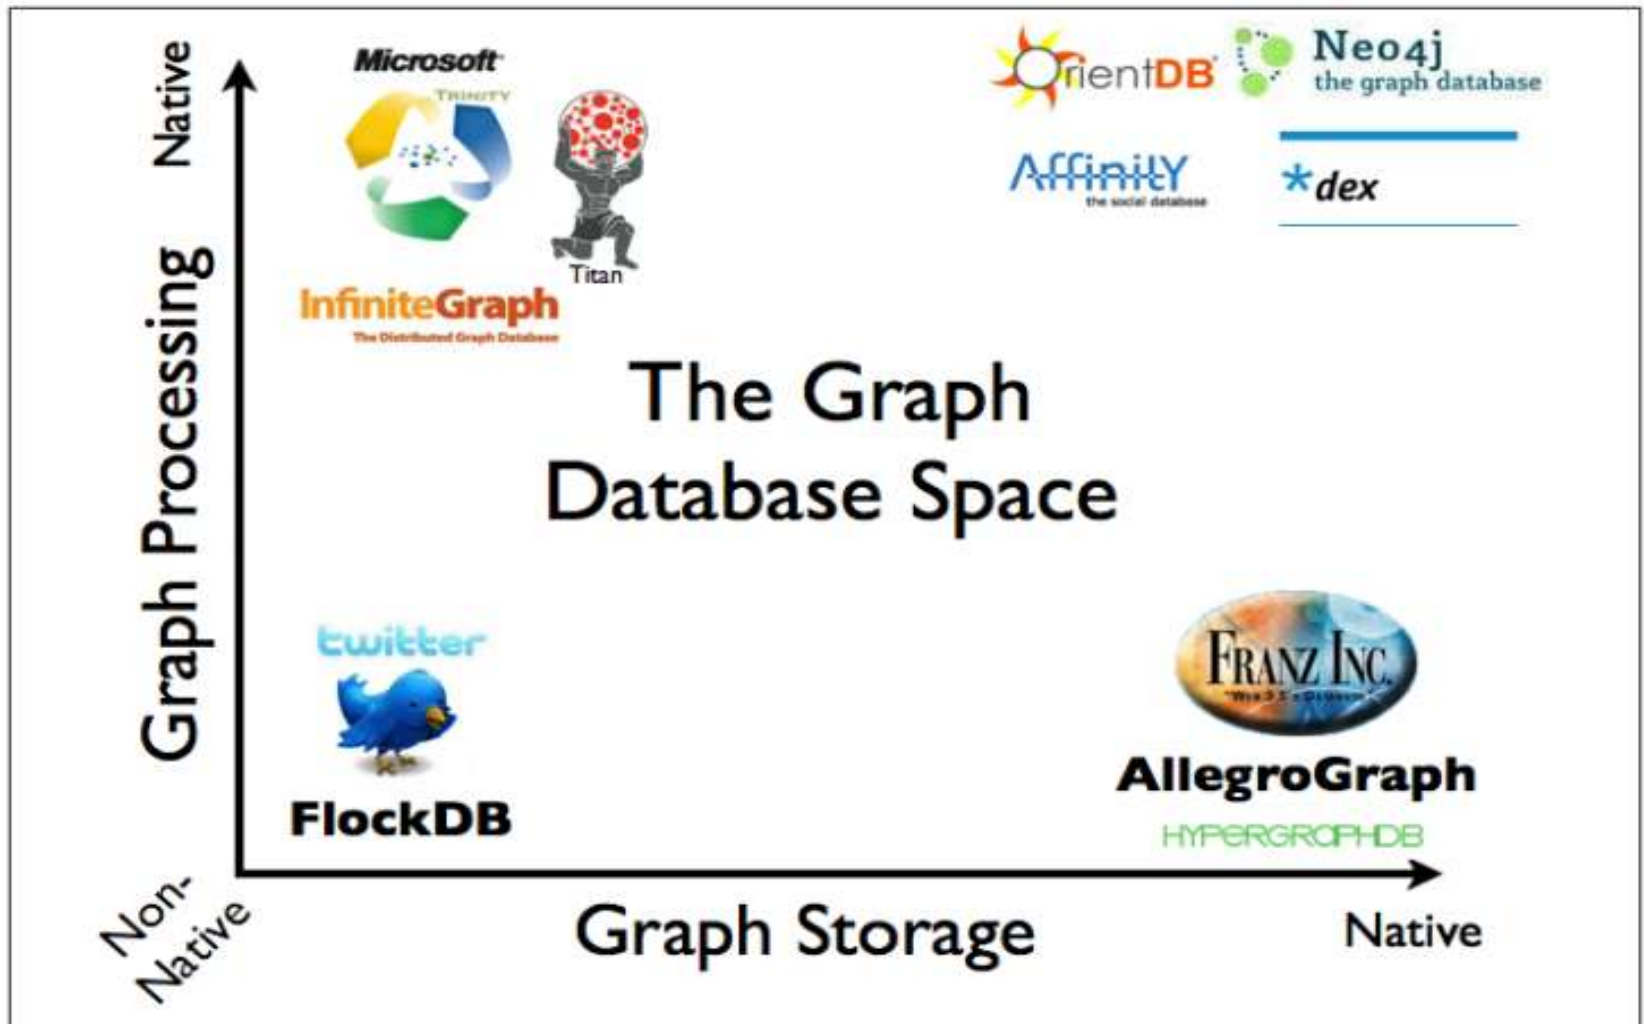
Main features: Neo4j

- ***intuitive***, using a graph model for data representation
- ***reliable***, with full ACID transactions
- ***durable and fast***, using a custom disk-based, native storage engine
- ***massively scalable***, up to several billion nodes/relationships/properties
- ***highly-available***, when distributed across multiple machines
- ***expressive***, with a powerful, human readable **graph query language**.
- ***fast***, with a powerful traversal framework for high-speed graph queries
- ***embeddable***, with a few small jars
- ***simple***, accessible by a convenient [REST interface](#) or an object-oriented Java [API](#)

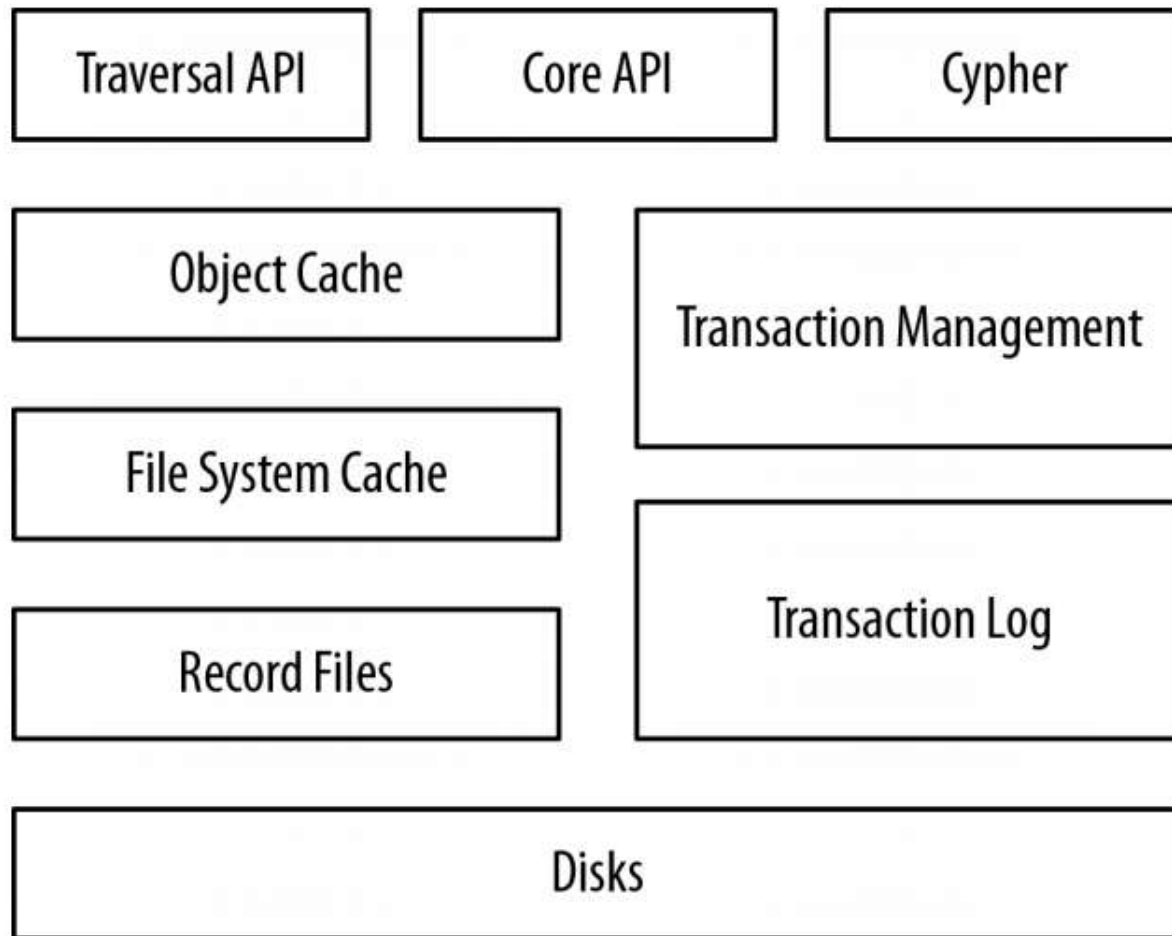
Neo4j in graph



An overview of graph database space:



Neo4j Architecture



Store files

- Neo4j stores graph data in a number of different store files.
- Each store file contains the data for a specific part of the graph (e.g., nodes, relationships, properties)
 - neostore.nodestore.db
 - neostore.relationshipstore.db
 - neostore.propertystore.db
 - neostore.propertystore.db.index
 - neostore.propertystore.db.strings
 - neostore.propertystore.db.arrays

Node store


- neostore.nodestore.db
- Size: 9 bytes
 - 1st byte: in-use flag
 - Next 4 bytes: ID of first relationship
 - Last 4 bytes: ID of first property of node
- Fixed size records enable fast lookups

Relationship store

- neostore.relationshipstore.db
- Size: 33 bytes
 - 1st byte: In use flag
 - Next 8 bytes: IDs of the nodes at the start and end of the relationship
 - 4 bytes: Pointer to the relationship type
 - 16 bytes: pointers for the next and previous relationship records for each of the start and end nodes. (property chain)
 - 4 bytes: next property id


Node/property record structure

inUse

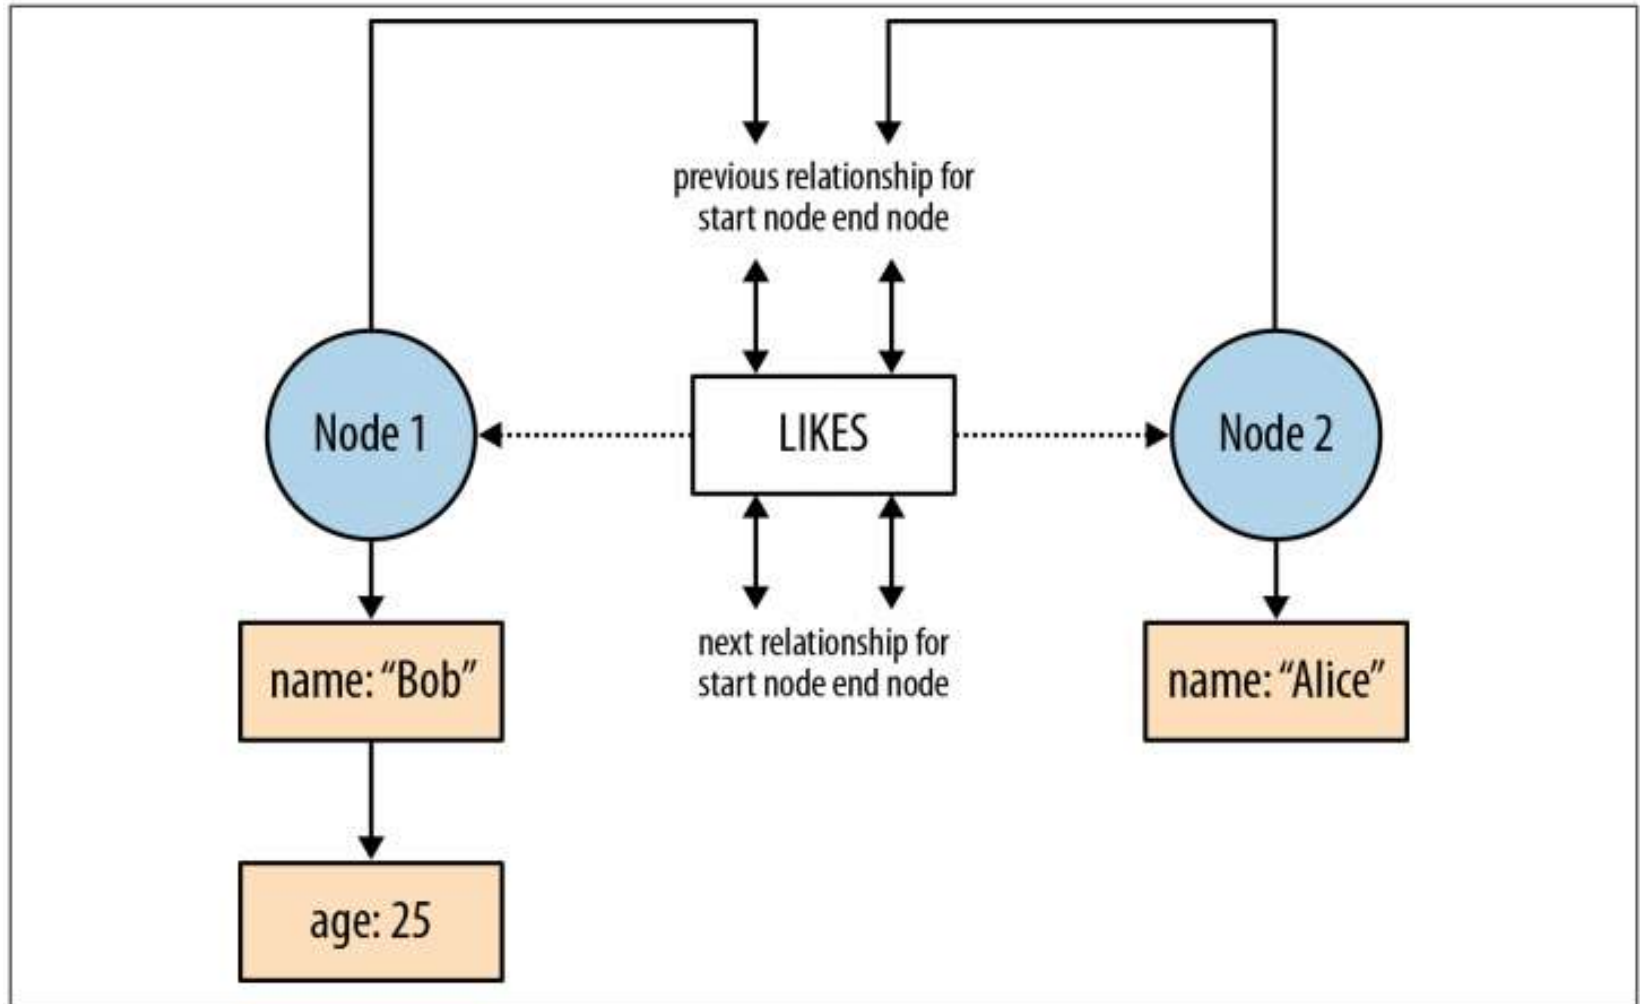


Relationship (33 bytes)

inUse



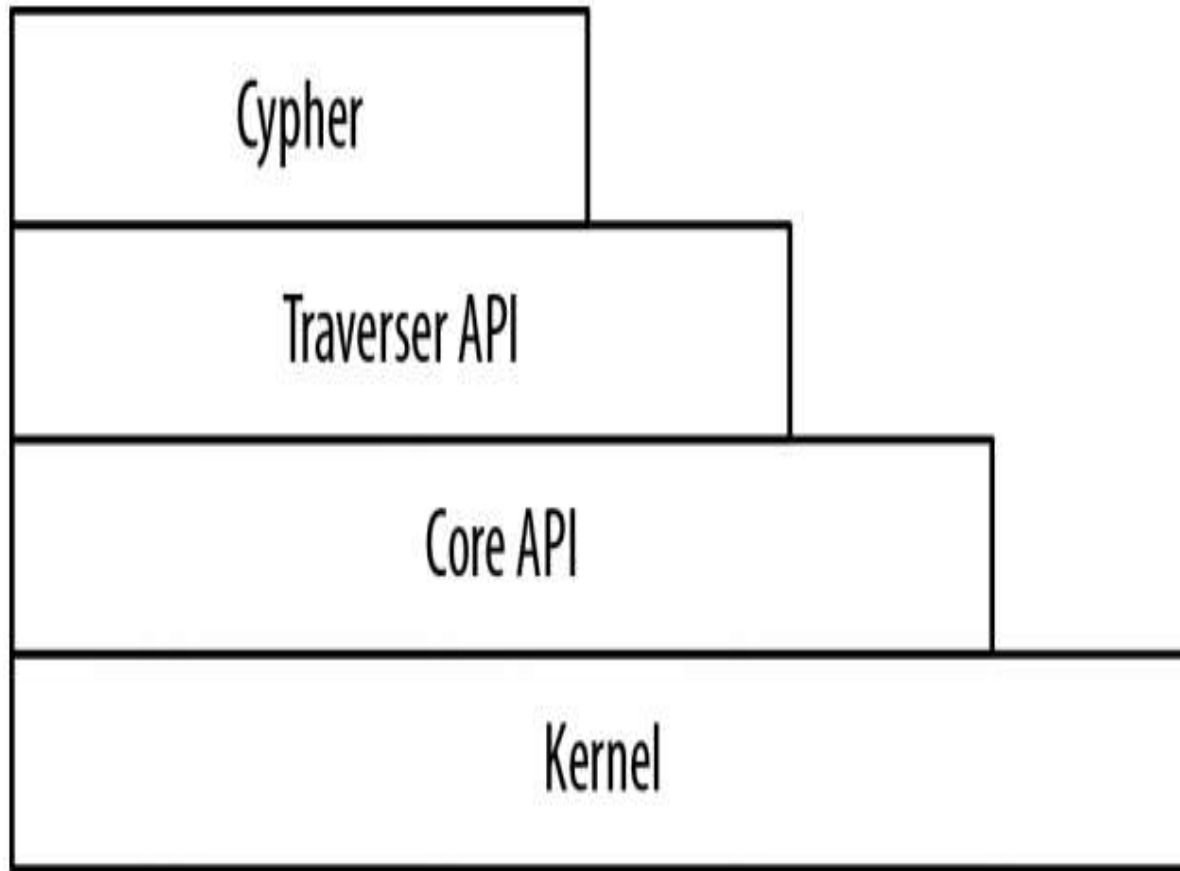
How a graph is physically stored



Neo4j: Data Size

nodes	2^{35} (~ 34 billion)
relationships	2^{35} (~ 34 billion)
properties	2^{36} to 2^{38} depending on property types (maximum ~ 274 billion, always at least ~ 68 billion)
relationship types	2^{15} (~ 32 000)

Logical view of neo4j user API




Application Architecture

- Embedded Neo4j
 - Low latency
 - Choice of APIs
 - Explicit transactions
 - JVM only
 - GC behavior
 - Database life cycle

Application Architecture (contd..)

- Server Mode
 - Using same embedded instance of neo4j
 - REST API
 - Platform independent
 - Isolation from application GC behavior
 - Network overhead
 - Support for server extensions


Benchmarking neo4j FOAF Query Possibilities by
Core API outperforms Cypher by about one order of magnitude in neo4j1.7
For newer neo4j versions Cypher becomes worse
by <http://www.rene-pickhardt.de>



FOAF like Queries per Second with various neo4j Query frame works

by <http://www.rene-pickhardt.de>

averages are built on experimants with more than 1k queries



Datasets

- Cineasts Movies & Actors
- Wordnet v:2.0
- Neo4j version 1.9 supports 10 of billions of nodes/relationships/properties
- The Neo4j team has publicly expressed the intention to support 100B+ nodes/relationships/properties in a single graph as part of its 2013 roadmap

benchmarks

- <https://code.google.com/p/orient/wiki/GraphDBComparison>
- <http://nosql.mypopescu.com/post/1451025794/neo4j-and-orientdb-performance-compared>
- <https://groups.google.com/forum/#!msg/orient-database/9J5s4q3WKxY/d5XrIpiVG6gJ>
- <https://baach.de/Members/jhb/neo4j-performance-compared-to-mysql>
- <http://www.slideshare.net/thobe/nosqleu-graph-databases-and-neo4j>
- <http://java.dzone.com/articles/get-full-neo4j-power-using>

Performance comparison: neo4j vs RDBMS (table 2-1 pg:20)

Supports index-free adjacency

Depth	RDBMS execution time (s)	Neo4j execution time (s)	Records returned
2	0.016	0.01	~2500
3	30.267	0.168	~110,000
4	1543.505	1.359	~600,000
5	Unfinished	2.132	~800,000
